

```
In [1]: import numpy as np
from tensorflow.keras.datasets import imdb
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Embedding, Flatten
from tensorflow.keras.preprocessing.sequence import pad_sequences

# Load the IMDB dataset
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=10000)

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz
17464789/17464789 [=====] - 0s 0us/step
```

```
In [ ]: max_len = 500

# Pad and truncate the sequences
x_train = pad_sequences(x_train, maxlen=max_len)
x_test = pad_sequences(x_test, maxlen=max_len)
```

```
In [ ]: model = Sequential()
model.add(Embedding(10000, 32, input_length=max_len))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(1, activation='sigmoid'))

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
In [ ]: model.fit(x_train, y_train, validation_split=0.2, epochs=5, batch_size=128)

Epoch 1/5
157/157 [=====] - 9s 45ms/step - loss: 0.5280 - accuracy: 0.6998 - val_loss: 0.3154 - val_accuracy: 0.8610
Epoch 2/5
157/157 [=====] - 7s 43ms/step - loss: 0.1789 - accuracy: 0.9330 - val_loss: 0.3099 - val_accuracy: 0.8704
Epoch 3/5
157/157 [=====] - 7s 44ms/step - loss: 0.0487 - accuracy: 0.9888 - val_loss: 0.3510 - val_accuracy: 0.8766
Epoch 4/5
157/157 [=====] - 7s 44ms/step - loss: 0.0122 - accuracy: 0.9990 - val_loss: 0.4307 - val_accuracy: 0.8722
Epoch 5/5
157/157 [=====] - 7s 44ms/step - loss: 0.0040 - accuracy: 0.9998 - val_loss: 0.4630 - val_accuracy: 0.8738
Out[ ]: <keras.callbacks.History at 0x2197f81d880>
```

```
In [ ]: loss, accuracy = model.evaluate(x_test, y_test)
print(f'Test accuracy: {accuracy * 100:.2f}%')

782/782 [=====] - 2s 3ms/step - loss: 0.4616 - accuracy: 0.8688
Test accuracy: 86.88%
```

```
In [ ]: def predict_review(review):
# Convert the review to a sequence of word indices
seq = imdb.get_word_index()
```

```

words = review.split()
seq = [seq[w] if w in seq else 0 for w in words]
seq = pad_sequences([seq], maxlen=max_len)

# Make the prediction
pred = model.predict(seq)[0]

# Return the prediction
return 'positive' if pred >= 0.5 else 'negative'

review = "This movie was great! I loved the story and the acting was superb."
prediction = predict_review(review)
print(f'Review: {review}')
print(f'Prediction: {prediction}')

```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imb_word_index.json
 1641221/1641221 [=====] - 2s 1us/step
 1/1 [=====] - 0s 62ms/step
 Review: This movie was great! I loved the story and the acting was superb.
 Prediction: positive

```

In [ ]: # Print model summary
        model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 500, 32)	320000
flatten (Flatten)	(None, 16000)	0
dense (Dense)	(None, 128)	2048128
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 1)	129

=====
 Total params: 2,368,257
 Trainable params: 2,368,257
 Non-trainable params: 0

```

In [ ]: from sklearn.metrics import confusion_matrix
        import matplotlib.pyplot as plt
        import numpy as np

        # Get predicted labels
        y_pred = np.round(model.predict(x_test))

        # Generate confusion matrix
        cm = confusion_matrix(y_test, y_pred)

        # Normalize confusion matrix
        cm_norm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

        # Set up plot
        fig, ax = plt.subplots(figsize=(8, 8))

```

```

# Plot confusion matrix
im = ax.imshow(cm_norm, interpolation='nearest', cmap=plt.cm.Reds)
ax.figure.colorbar(im, ax=ax)

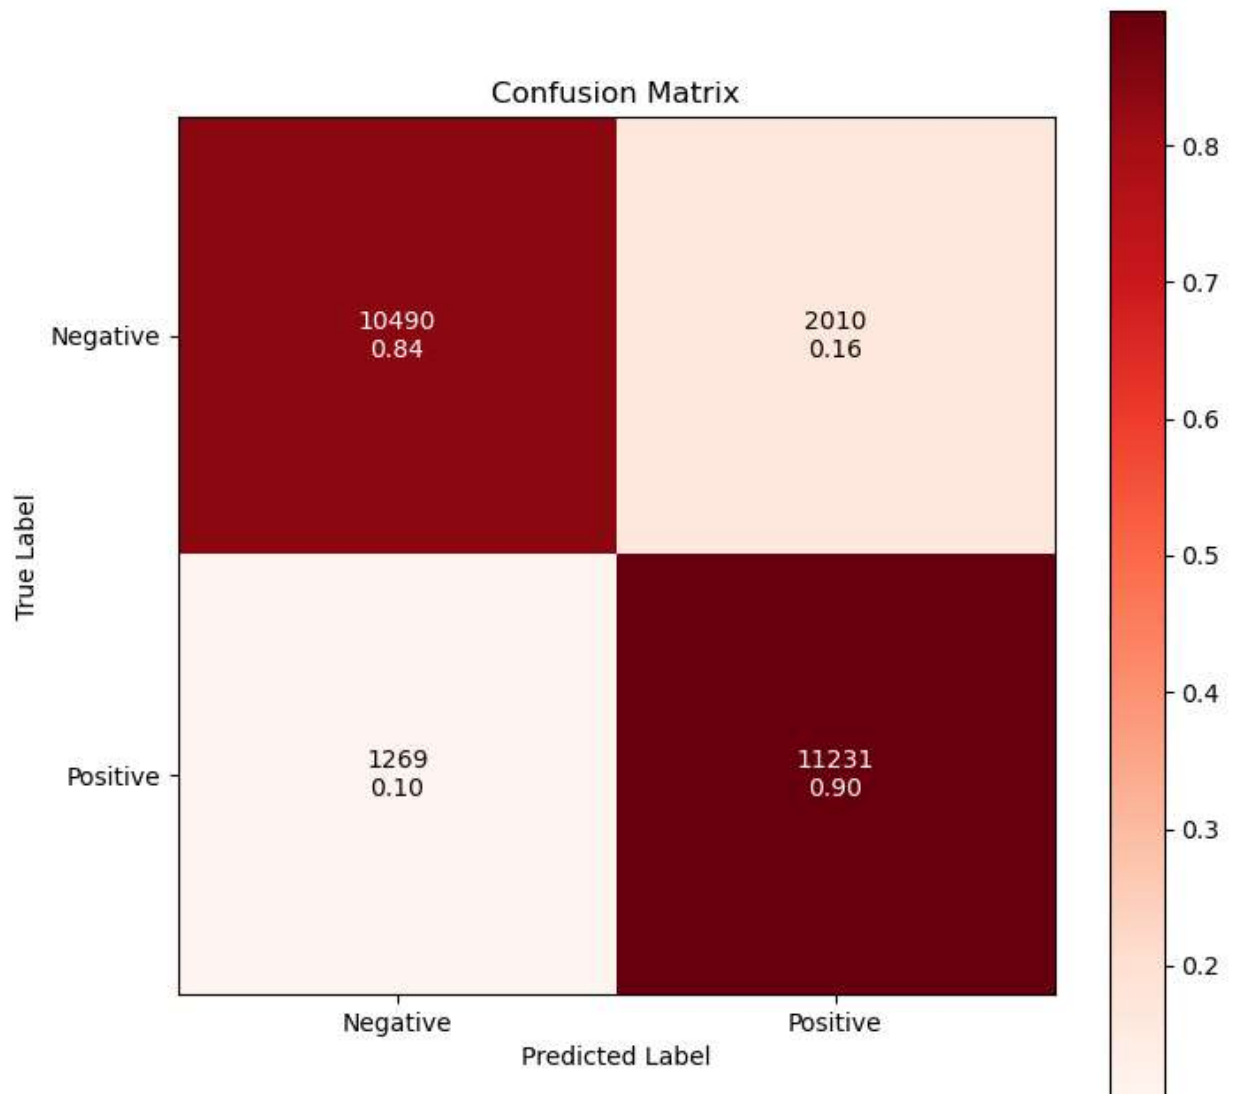
# Set Labels
ax.set(xticks=np.arange(cm.shape[1]),
       yticks=np.arange(cm.shape[0]),
       xticklabels=['Negative', 'Positive'], yticklabels=['Negative', 'Positive'],
       title='Confusion Matrix',
       ylabel='True Label',
       xlabel='Predicted Label')

# Add Labels to each cell
thresh = cm_norm.max() / 2.
for i in range(cm_norm.shape[0]):
    for j in range(cm_norm.shape[1]):
        ax.text(j, i, format(cm[i, j], 'd') + '\n' + format(cm_norm[i, j], '.2f'),
                ha="center", va="center",
                color="white" if cm_norm[i, j] > thresh else "black")

# Show plot
plt.show()

```

782/782 [=====] - 2s 3ms/step



```
In [ ]: from sklearn.metrics import classification_report  
  
print(classification_report(y_test, y_pred, target_names=['Negative', 'Positive']))
```

	precision	recall	f1-score	support
Negative	0.89	0.84	0.86	12500
Positive	0.85	0.90	0.87	12500
accuracy			0.87	25000
macro avg	0.87	0.87	0.87	25000
weighted avg	0.87	0.87	0.87	25000

```
In [ ]:
```